

Open-Source Report

Proof of knowing your stuff in CSE312

Guidelines

Provided below is a template you must use to write your reports for your project.

Here are some things to note when working on your report, specifically about the **General Information & Licensing** section for each technology.

- **Code Repository:** Please link the code and not the documentation. If you'd like to refer to the documentation in the **Magic** section, you're more than welcome to, but we need to see the code you're referring to as well.
- **License Type:** Three letter acronym is fine.
- **License Description:** No need for the entire license here, just what separates it from the rest.
- **License Restrictions:** What can you *not* do as a result of using this technology in your project? Some licenses prevent you from using the project for commercial use, for example.

Also, feel free to extend the cell of any section if you feel you need more room.

If there's anything we can clarify, please don't hesitate to reach out! You can reach us using the methods outlined on the course website or see us during our office hours.


FastAPI

General Information & Licensing

Code Repository	https://github.com/tiangolo/fastapi
License Type	MIT License
License Description	<ul style="list-style-type: none">• Permissive license• Allows unrestricted use without any limitation free of charge• Protects the copyright holders from liability for any consequences caused by anything utilizing the licensed code
License Restrictions	<ul style="list-style-type: none">• Derivative code must also be licensed under the MIT License

FastAPI is a Python web framework that allows for rapid development of REST APIs. It is built on top of Starlette, which is a lightweight ASGI framework. FastAPI and Starlette do not ship with server code, but rather rely on an additional ASGI webserver to handle network connections and pass data to the Python application following the ASGI specification. As such, understanding the TCP connection requires understanding the ASGI webserver.

For our implementation, we used the ASGI webserver uvicorn. This is the program that actually opens sockets and handles network connections, so it will be the focus of the following discussion. At a high level, uvicorn opens a socket by invoking an API in the Python standard library asyncio. Asyncio then handles connection requests, reads data off of the TCP socket, and sends this data back to uvicorn via an asyncio protocol. Uvicorn then passes the raw bytes to the Python library h11 (which stands for HTTP 1.1). This library parses the bytes into a request object, which uvicorn then packs into an ASGI message and sends to FastAPI and Starlette. FastAPI and Starlette then generate a request object from the ASGI message, pass it through a series of middleware, and finally pass it to the application code. The application code generates the response, the response is returned to FastAPI and Starlette, and the process repeats in reverse. A diagram of the information flow is shown below.



Since we did our development on Windows, we will be describing the chain of calls that use the Windows APIs. Specifically, the library used to interact with the TCP socket, asyncio, uses different Python modules to manage sockets depending on the platform. Since we are using Windows with CPython, the modules of interest are `_winapi` and `_overlapped`. The functions in these modules perform the system calls that are necessary to perform the TCP socket IO for Windows devices. The full chain of calls is described below.

First, we call `uvicorn.run()` on line 144 in `main.py` to start our ASGI webserver
<https://github.com/mattrubino/cse312-group-project/blob/main/api/src/main.py#L144>

```
142
143     if __name__ == "__main__":
144         uvicorn.run("main:app", host="0.0.0.0", reload=DEV, proxy_headers=True)
```

`uvicorn.run()` then initializes a config, instantiates a `Server` object, and then calls the `Server` class's `run()` function on line 578 in `uvicorn's main.py`
<https://github.com/encode/uvicorn/blob/master/uvicorn/main.py#L578>

```
576         Multiprocess(config, target=server.run, sockets=[sock]).run()
577     else:
578         server.run()
579         if config.uds and os.path.exists(config.uds):
580             os.remove(config.uds) # pragma: py-win32
```

The `Server's run()` method then runs the `Server's serve()` method in a new `asyncio` event loop on line 61 of `server.py`. This starts the asynchronous processing, which will help in achieving high throughput
<https://github.com/encode/uvicorn/blob/master/uvicorn/server.py#L61>

```
59     def run(self, sockets: Optional[List[socket.socket]] = None) -> None:
60         self.config.setup_event_loop()
61         return asyncio.run(self.serve(sockets=sockets))
```

The `Server's serve()` method loads the config, logs some information, and then yields control to its `startup()` method on line 78 of `server.py`
<https://github.com/encode/uvicorn/blob/master/uvicorn/server.py#L78>

```
76         logger.info(message, process_id, extra={"color_message": message})
77
78         await self.startup(sockets=sockets)
79         if self.should_exit:
80             return
```

The `startup` method defines a factory function `create_protocol()` on line 96. This factory function takes in an `asyncio` event loop and returns an `asyncio` protocol object. `Asyncio` will use this factory function to generate a protocol object. This protocol object is the mechanism by which `asyncio` will eventually pass socket connection information and socket data back to `uvicorn` (more on this later). Finally, note that the default protocol specified in the configuration is the `H11Protocol` (standing for `HTTP 1.1`), which is defined in `uvicorn/protocols/http/h11_impl.py`
<https://github.com/encode/uvicorn/blob/master/uvicorn/server.py#L96>
https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L78

```

94         config = self.config
95
96     def create_protocol(
97         _loop: Optional[asyncio.AbstractEventLoop] = None,
98     ) -> asyncio.Protocol:
99         return config.http_protocol_class( # type: ignore[call-arg]
100             config=config,
101             server_state=self.server_state,
102             app_state=self.lifespan.state,
103             _loop=_loop,
104         )

```

```

78 class H11Protocol(asyncio.Protocol):
79     def __init__(
80         self,
81         config: Config,
82         server_state: ServerState,
83         app_state: Dict[str, Any],
84         _loop: Optional[asyncio.AbstractEventLoop] = None,
85     ) -> None:

```

Back to the Server's startup() method, this routine invokes the asyncio event loop method create_server() on line 161, passing the protocol factory function described above. This is how uvicorn offloads the socket management to asyncio

<https://github.com/encode/uvicorn/blob/master/uvicorn/server.py#L161>

```

160         try:
161             server = await loop.create_server(
162                 create_protocol,
163                 host=config.host,
164                 port=config.port,
165                 ssl=config.ssl,
166                 backlog=config.backlog,
167             )

```

We will now look at the source code of asyncio, since this is where the socket management occurs

The create_server() method is defined on line 1444 of asyncio/base_events.py. Critically, this method creates an instance of the Python socket object using the information specified in the configuration on line 1513

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1513

```

1509         try:
1510             for res in infos:
1511                 af, socktype, proto, canonname, sa = res
1512                 try:
1513                     sock = socket.socket(af, socktype, proto)
1514                     except socket.error:

```

This socket object is then bound to the IP and port based on the configuration information on line 1537. This is necessary so that the operating system can verify that the specified address and port are unused

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1537

```

1536         try:
1537             sock.bind(sa)
1538         except OSError as err:
1539             raise OSError(err.errno, 'error while attempting '
1540                           'to bind on address %r: %s'

```

This socket object is then passed to an instance of the asyncio Server class, and the `_start_serving()` method is called on the server object on line 1561

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1561

```

1560         if start_serving:
1561             server._start_serving()
1562             # Skip one loop iteration so that all

```

The `_start_serving()` method starts the listening socket on line 315, which is necessary to start accepting connections. It then runs the `_start_serving()` method on the loop object on line 316

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L316

```

314         for sock in self._sockets:
315             sock.listen(self._backlog)
316             self._loop._start_serving(
317                 self._protocol_factory, sock, self._ssl_context,
318                 self, self._backlog, self._ssl_handshake_timeout,
319                 self._ssl_shutdown_timeout)

```

The loop object is an instance of the `BaseProactorEventLoop` class. This class manages an asyncio event loop. The `_start_serving()` method of the `BaseProactorEventLoop` is defined on line 836 of `proactor_events.py`, which schedules the nested function `loop()` to run on the next event loop cycle on line 879 using the `call_soon()` routine

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L879

```

878
879         self.call_soon(loop)
880

```

In the next event loop cycle, the inner loop() function runs. In this function, the accept() method of the proactor object is called on line 861 and assigned to the variable f. This function returns a future which completes when a new TCP connection is accepted. When this occurs, the loop function will run again. This is because the function is scheduled as a done callback of the future on line 877, effectively turning the function into an infinite loop
https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L861
https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L877

```
859         if self.is_closed():
860             return
861         f = self._proactor.accept(sock)
862         except OSError as exc:
863             if sock.fileno() != -1:
864                 self.call_exception_handler({
875                     else:
876                         self._accept_futures[sock.fileno()] = f
877                 f.add_done_callback(loop)
878
```

Inside the future, the function waits for connection requests. On Windows, the proactor object mentioned above is an instance of ProactorEventLoop (defined in windows_events.py). The accept() method on this class is defined on line 543. This method retrieves the information of the remote socket using the _get_accept_socket() method on line 545. It then accepts the TCP connection using the AcceptEx function from the _overlapped module (this comes directly from the [Windows API](#), it is invoked using a Python C extension module) on line 547
https://github.com/python/cpython/blob/main/Lib/asyncio/windows_events.py#L547

```
543     def accept(self, listener):
544         self._register_with_ioep(listener)
545         conn = self._get_accept_socket(listener.family)
546         ov = _overlapped.Overlapped(NULL)
547         ov.AcceptEx(listener.fileno(), conn.fileno())
```

When the future completes, the loop() function is executed as a callback. This time, the variable f holds the future that just completed, so the connection and address can be retrieved on line 844
https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L844

```
843         if f is not None:
844             conn, addr = f.result()
845             if self._debug:
846                 logger.debug("%r got a new connection from %r: %r",
```

The socket transport is then created on line 856, passing the connection information and the uvicorn-supplied protocol implementation
https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L856

```

855         else:
856             self._make_socket_transport(
857                 conn, protocol,
858                 extra={'peername': addr}, server=server)

```

The `_make_socket_transport()` function creates a new instance of `_ProactorSocketTransport` and returns it on line 647

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L647

```

645     def _make_socket_transport(self, sock, protocol, waiter=None,
646                               extra=None, server=None):
647         return _ProactorSocketTransport(self, sock, protocol, waiter,
648                                         extra, server)

```

Inside the constructor of `_ProactorSocketTransport`, something critical happens. Specifically, `_ProactorSocketTransport` extends `_ProactorReadPipeTransport`, which extends `_ProactorBasePipeTransport`, and something important happens in the constructor of `_ProactorBasePipeTransport`. On line 67, the `connection_made()` function of the protocol is scheduled to run in the event loop, putting it inside the loop's queue `_ready`. Specifically, this will be the function `H11Protocol.connection_made()`. This function is the mechanism by which `asyncio` passes connection information back to `uvicorn`

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L67

```

65         if self._server is not None:
66             self._server._attach()
67         self._loop.call_soon(self._protocol.connection_made, self)
68         if waiter is not None:
69             # only wake up the waiter when connection_made() has been called

```

After `connection_made()` is scheduled, the function `_loop_reading()` is also scheduled with the event loop, similarly putting it inside the `_ready` queue. This occurs on line 192 in the constructor of `_ProactorReadPipeTransport`. Importantly, this routine contains the code for reading bytes off of the TCP connection

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L192

```

191         self._data = bytearray(buffer_size)
192         self._loop.call_soon(self._loop_reading)
193         self._paused = False
194

```

As the event loop runs, it will first pop the handle for the `H11Protocol.connection_made` function off of the `_ready` queue on line 1936

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1936

```

1934         ntodo = len(self._ready)
1935         for i in range(ntodo):
1936             handle = self._ready.popleft()
1937             if handle._cancelled:
1938                 continue

```

It then runs the handle on line 1951

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1951

```

1950         else:
1951             handle._run()
1952             handle = None # Needed to break cycles when an exception occurs.

```

This triggers the connection handling code on uvicorn on line 128 of uvicorn/protocols/http/h11_impl.py, which registers the connection with uvicorn

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L128

```

127         # Protocol interface
128         def connection_made( # type: ignore[override]
129             self, transport: asyncio.Transport
130         ) -> None:
131             self.connections.add(self)

```

On the next tick of the event loop, it pops off the handle for

`_ProactorReadPipeTransport._loop_reading()` off of the `_ready` queue on line 1936

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1936

```

1934         ntodo = len(self._ready)
1935         for i in range(ntodo):
1936             handle = self._ready.popleft()
1937             if handle._cancelled:
1938                 continue

```

It then runs the handle on line 1951

https://github.com/python/cpython/blob/main/Lib/asyncio/base_events.py#L1951

```

1950         else:
1951             handle._run()
1952             handle = None # Needed to break cycles when an exception occurs.

```

This triggers the reading of data from the TCP socket. Specifically, it triggers the `_loop_reading()` method on line 276 of `asyncio/proactor_events.py`. This calls the method `recv_into` on the `_proactor` object on line 306, which reads data from the supplied socket

into the supplied buffer `_data`

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L306

```
304         if not self._paused:
305             # reschedule a new read
306             self._read_fut = self._loop._proactor.recv_into(self._sock, self._data)
307         except ConnectionAbortedError as exc:
308             if not self._closing:
```

The `_proactor` object is an instance of `loopProactor`, so this triggers the `receive_into()` method on line 482 of `asyncio/windows_events.py`. This runs the `WSARecvInto()` method from the `_overlapped` module on line 487. This again invokes the [Windows API](#), but this time for reading data from the socket

https://github.com/python/cpython/blob/main/Lib/asyncio/windows_events.py#L487

```
485         try:
486             if isinstance(conn, socket.socket):
487                 ov.WSARecvInto(conn.fileno(), buf, flags)
488             else:
489                 ov.ReadFileInto(conn.fileno(), buf)
```

After reading some data into the buffer, the `_loop_reading()` method reschedules itself on line 322. This is necessary in case the first read did not read all of the data. This implements the buffering of TCP data recursively

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L322

```
320         else:
321             if not self._paused:
322                 self._read_fut.add_done_callback(self._loop_reading)
323             finally:
```

The data in the buffer is then passed to the `_data_received()` method in the `finally` block on line 325

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L325

```
323         finally:
324             if length > -1:
325                 self._data_received(data, length)
```

The `_data_received()` method performs some checks of the supplied inputs and then passes them to the `H11Protocol` instance from `uvicorn` on line 274

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L274

```
272         return
273     else:
274         self._protocol.data_received(data)
275
```

Inside the `data_receive()` method of `uvicorn`, the data is passed into the `receive_data()` method of the connection object on line 201. This attaches the received data buffer to the connection

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L201

```
198         def data_received(self, data: bytes) -> None:
199             self._unset_keepalive_if_required()
200
201             self.conn.receive_data(data)
202             self.handle_events()
```

On the next line of the `data_receive()` method, `handle_events()` is called, which will process the data attached to the connection in the previous step

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L202

```
198         def data_received(self, data: bytes) -> None:
199             self._unset_keepalive_if_required()
200
201             self.conn.receive_data(data)
202             self.handle_events()
```

Inside of `handle_events()`, the previous event data is retrieved using the method `next_event()` on line 207. The details of this method are described below

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L207

```
204         def handle_events(self) -> None:
205             while True:
206                 try:
207                     event = self.conn.next_event()
208                     except h11.RemoteProtocolError:
```

Inside `next_event()`, the method `_extract_next_receive_event()` is called in `h11/_connection.py` on line 469

https://github.com/python-hyper/h11/blob/master/h11/_connection.py#L469

```
468         try:
469             event = self._extract_next_receive_event()
470             if event not in [NEED_DATA, PAUSED]:
471                 self._process_event(self.their_role, cast(Event, event))
```

Inside `_extract_next_receive_event()`, the data in the receive buffer is passed to the `_reader()` function on line 411

https://github.com/python-hyper/h11/blob/master/h11/_connection.py#L411

```

410         assert self._reader is not None
411         event = self._reader(self._receive_buffer)
412         if event is None:

```

The `_reader` identifier is a reference to the function `_maybe_read_from_IDLE_client()` on line 75 of `h11/_readers.py`. This function then calls the method `maybe_extract_lines()` on the buffer on line 76

https://github.com/python-hyper/h11/blob/master/h11/_readers.py#L76

```

75     def maybe_read_from_IDLE_client(buf: ReceiveBuffer) -> Optional[Request]:
76         lines = buf.maybe_extract_lines()
77         if lines is None:
78             if buf.is_next_line_obviously_invalid_request_line():

```

The `maybe_extract_lines()` method parses the byte buffer into a list of bytearrays, where each bytearray represents a line of the HTTP request. It does this by performing a split on `"\n"` on line 126 in `h11/_receivebuffer.py` (and then removing trailing `"\r"`s)

https://github.com/python-hyper/h11/blob/master/h11/_receivebuffer.py#L126

```

123         # Truncate the buffer and return it.
124         idx = match.span(0)[-1]
125         out = self._extract(idx)
126         lines = out.split(b"\n")
127
128         for line in lines:
129             if line.endswith(b"\r"):
130                 del line[-1]

```

The HTTP request line is then validated on line 411 of `h11/_readers.py` because you can never trust your users!

https://github.com/python-hyper/h11/blob/master/h11/_readers.py#L83

```

82         raise LocalProtocolError("no request line received")
83         matches = validate(
84             request_line_re, lines[0], "illegal request line: {!r}", lines[0]
85         )

```

Finally, the reader parses the headers (described in detail in the header report) and returns a `Request` object that wraps the parsed data on line 86. This object is then returned all the way back up to `handle_events()`

https://github.com/python-hyper/h11/blob/master/h11/_readers.py#L86

```

86         return Request(
87             headers=list(_decode_header_lines(lines[1:])), _parsed=True, **matches
88         )

```

The event has type Request, so the conditional on line 226 of `uvicorn/protocols/http/h11_impl.py` evaluates to true. In this branch, an ASGI scope dictionary is created on line 229. This dictionary follows the ASGI specification, and is necessary for passing a message to an ASGI application like Starlette

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L229

```
228         raw_path, _, query_string = event.target.partition(b"?.")
229         self.scope = { # type: ignore[typeddict-item]
230             "type": "http",
231             "asgi": {
232                 "version": self.config.asgi_version,
233                 "spec_version": "2.3",
234             },
235             "http_version": event.http_version.decode("ascii"),
236             "server": self.server,
```

This scope object is passed into an instance of `RequestResponseCycle` created on line 264, and the ASGI request is scheduled to run in the event loop on line 276

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L276

```
275         )
276         task = self.loop.create_task(self.cycle.run_asgi(app))
277         task.add_done_callback(self.tasks.discard)
278         self.tasks.add(task)
```

Inside the `run_asgi()` method, the ASGI app (Starlette/FastAPI) is invoked on line 428 according to the ASGI specification

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L428

```
426     async def run_asgi(self, app: "ASGI3Application") -> None:
427         try:
428             result = await app( # type: ignore[func-returns-value]
429                 self.scope, self.receive, self.send
430             )
```

Uvicorn automatically wraps the user's application in middleware, so this middleware gets invoked first. Specifically, the invocation mentioned previously triggers the `__call__` method of the `ProxyHeadersMiddleware` on line 49, which automatically manages header information for proxies

https://github.com/encode/uvicorn/blob/master/uvicorn/middleware/proxy_headers.py#L49

```
48
49     async def __call__(
50         self, scope: "Scope", receive: "ASGIReceiveCallable", send: "ASGISendCallable"
51     ) -> None:
52         if scope["type"] in ("http", "websocket"):
53             scope = cast(Union["HTTPScope", "WebSocketScope"], scope)
54             client_addr: Optional[Tuple[str, int]] = scope.get("client")
55             client_host = client_addr[0] if client_addr else None
56
```

This middleware then invokes the FastAPI application on line 78

https://github.com/encode/uvicorn/blob/master/uvicorn/middleware/proxy_headers.py#L78

```
76         scope["client"] = (host, port) # type: ignore[arg-type]
77
78         return await self.app(scope, receive, send)
```

The FastAPI application entrypoint is the `__call__` method of the FastAPI type, defined on line 273 of `fastapi/applications.py`. This method is also a form of middleware, simply passing the ASGI message to the superclass Starlette on line 276

<https://github.com/tiangolo/fastapi/blob/master/fastapi/applications.py#L276>

```
273     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
274         if self.root_path:
275             scope["root_path"] = self.root_path
276         await super().__call__(scope, receive, send)
```

The Starlette application entrypoint is the `__call__` method of the Starlette type, defined on line 118 of `starlette/applications.py`. This method assembles the middleware stack if it does not exist and then invokes it on line 122

<https://github.com/encode/starlette/blob/master/starlette/applications.py#L122>

```
118     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
119         scope["app"] = self
120         if self.middleware_stack is None:
121             self.middleware_stack = self.build_middleware_stack()
122         await self.middleware_stack(scope, receive, send)
```

The ASGI request data must then be passed through Starlette's entire middleware stack. Specifically, it starts by going through `ServerErrorMiddleware` in `errors.py`, then `ExceptionMiddleware` in `exceptions.py`, then `AsyncExitStackMiddleware` in `asyncexitstack.py`, and then it finally reaches `Router` in `routing.py`

<https://github.com/encode/starlette/blob/master/starlette/middleware/errors.py#L147>

<https://github.com/encode/starlette/blob/master/starlette/middleware/exceptions.py#L53>

<https://github.com/tiangolo/fastapi/blob/master/fastapi/middleware/asyncexitstack.py#L12>

<https://github.com/encode/starlette/blob/master/starlette/routing.py#L697>

```
147     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
148         if scope["type"] != "http":
149             await self.app(scope, receive, send)
150             return
151
```

```
53     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
54         if scope["type"] not in ("http", "websocket"):
55             await self.app(scope, receive, send)
56             return
57
```

```
12     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
13         if AsyncExitStack:
14             dependency_exception: Optional[Exception] = None
15             async with AsyncExitStack() as stack:
16                 scope[self.context_name] = stack
```

```

697     async def __call__(self, scope: Scope, receive: Receive, send: Send) -> None:
698         """
699         The main entry point to the Router class.
700         """
701         assert scope["type"] in ("http", "websocket", "lifespan")

```

Inside the router middleware, a loop over all the routes is performed on line 712. When a route that matches the path is found, the corresponding handle method is called on line 718

<https://github.com/encode/starlette/blob/master/starlette/routing.py#L718>

```

712         for route in self.routes:
713             # Determine if any route matches the incoming scope,
714             # and hand over to the matching route if found.
715             match, child_scope = route.matches(scope)
716             if match == Match.FULL:
717                 scope.update(child_scope)
718             await route.handle(scope, receive, send)
719         return

```

The handle method then runs the supplied ASGI app on line 276

<https://github.com/encode/starlette/blob/master/starlette/routing.py#L276>

```

274             await response(scope, receive, send)
275         else:
276             await self.app(scope, receive, send)

```

This triggers the app closure defined within the request_response method on line 63 of starlette/routing.py. This app function passes the request object to the handler function on line 66.

<https://github.com/encode/starlette/blob/master/starlette/routing.py#L66>

```

63     async def app(scope: Scope, receive: Receive, send: Send) -> None:
64         request = Request(scope, receive=receive, send=send)
65         if is_coroutine:
66             response = await func(request)

```

This triggers the app closure defined within the get_request_handler method on line 168 of fastapi/routing.py. The app function does some additional parsing and then passes the request data to run_endpoint_function() on line 237

<https://github.com/encode/fastapi/blob/master/fastapi/routing.py#L237>

```

236         else:
237             raw_response = await run_endpoint_function(
238                 dependant=dependant, values=values, is_coroutine=is_coroutine
239             )
240

```

In `run_endpoint_function()`, the request data is passed to the application through `dependent.call()` on line 163

<https://github.com/tiangolo/fastapi/blob/master/fastapi/routing.py#L163>

```
162         if is_coroutine:
163             return await dependant.call(**values)
```

At this point, the GET request for `/user/Matt` reaches the application code on line 77 of `main.py`

<https://github.com/matttrubino/cse312-group-project/blob/main/api/src/main.py#L77>

```
76 @app.get("/user/{username}")
77 async def user(response: Response, username: str) -> dict:
78     # Cannot look up invalid username
79     if not validUsername(username):
80         response.status_code = 400
81         return "Bad Request"
```

The `user()` function from our code then returns a dictionary of user data on line 92

<https://github.com/matttrubino/cse312-group-project/blob/main/api/src/main.py#L92>

```
90         del user["Password"]
91
92         return user
```

The raw response that is returned is then serialized in the call to `serialize_response()` on line 255 of `fastapi/routing.py`

<https://github.com/tiangolo/fastapi/blob/master/fastapi/routing.py#L255>

```
253         if sub_response.status_code:
254             response_args["status_code"] = sub_response.status_code
255         content = await serialize_response(
256             field=response_field,
257             response_content=raw_response,
258             include=response_model_include,
```

In `serialize_response()`, the dictionary is passed into `jsonable_encoder()` on line 142 and then returned

<https://github.com/tiangolo/fastapi/blob/master/fastapi/routing.py#L142>


```

140         if errors:
141             raise ValidationError(errors, field.type_)
142         return jsonable_encoder(
143             value,
144             include=include,
145             exclude=exclude,
146             by_alias=by_alias,
147             exclude_unset=exclude_unset,

```

In the `jsonable_encoder()` function, the type of the response value is inspected, and the appropriate code to pack it into a dictionary object is executed. In this case, the response value is a dictionary, so the conditional on line 92 of `fastapi/encoders.py` returns true. Inside this function, each key and value is passed into the `jsonable_encoder()` recursively on lines 109 and 117 respectively. Once the original call returns, the value will be a JSON-serializable Python dictionary

<https://github.com/tiangolo/fastapi/blob/master/fastapi/encoders.py#L109>

<https://github.com/tiangolo/fastapi/blob/master/fastapi/encoders.py#L117>

```

109         encoded_key = jsonable_encoder(
110             key,
111             by_alias=by_alias,
112             exclude_unset=exclude_unset,
113         )
114     encoded_value = jsonable_encoder(
115         value,
116         by_alias=by_alias,
117         exclude_unset=exclude_unset,

```

This returned dictionary is then passed to the identifier `actual_response_class` on line 266, which points to the constructor for `JSONResponse` in `starlette/responses.py` on line 196

<https://github.com/tiangolo/fastapi/blob/master/fastapi/routing.py#L266>

<https://github.com/encode/starlette/blob/master/starlette/responses.py#L196>

```

266         response = actual_response_class(content, **response_args)
267         if not is_body_allowed_for_status_code(response.status_code):
268             response.body = b""
269         response.headers.raw.extend(sub_response.headers.raw)
270     ) -> None:
271         super().__init__(content, status_code, headers, media_type, background)

```


The returned JSONResponse object is returned back up through the middleware to the app closure defined in request_response() mentioned previously. The response is then awaited on line 69 of starlette/routing.py, triggering the __call__ method of the Response class

<https://github.com/encode/starlette/blob/master/starlette/routing.py#L69>

```
67         else:
68             response = await run_in_threadpool(func, request)
69         await response(scope, receive, send)
70
71     return app
```

In the __call__ method of Response, two ASGI messages are sent. The first on line 164 signifies the start of an HTTP response with the status code and headers,

<https://github.com/encode/starlette/blob/master/starlette/responses.py#L164>

```
163     async def call_(self, scope: Scope, receive: Receive, send: Send) -> None:
164         await send(
165             {
166                 "type": "http.response.start",
167                 "status": self.status_code,
168                 "headers": self.raw_headers,
169             }
170         )
```

This message is passed up the middleware to the send() method defined in uvicorn/protocols/http/h11_impl.py on RequestResponseCycle. Inside this method, the response data is passed into an h11 Response object on line 509

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L509

```
508         reason = STATUS_PHRASES[status_code]
509         event = h11.Response(
510             status_code=status_code, headers=headers, reason=reason
511         )
```

This Response object is then passed to the send method on the connection object on line 512. This method from h11._connection.py converts the HTTP event into bytes that can be sent on the TCP connection

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L512

```
512         output = self.conn.send(event)
513         self.transport.write(output)
```

The output bytes are then passed to the write method on the transport object on line 513, which will send the raw bytes over the TCP socket

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L513

```
512         output = self.conn.send(event)
513         self.transport.write(output)
```

The transport object is an instance of `_ProactorBaseWritePipeTransport`. The write method of this class is defined on line 338 of `asyncio/proactor_events.py`. Importantly, on line 366, it invokes the method `_loop_writing`, passing a copy of the raw bytes

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L366

```
365         # Pass a copy, except if it's already immutable.
366         self._loop_writing(data=bytes(data))
367     elif not self._buffer: # WRITING -> BACKED UP
```

The first time the `_loop_writing` method runs, the `else` block will run. Importantly, this initializes the `_write_fut` identifier to be the future returned by calling `send` on the proactor object on line 402

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L402

```
401         else:
402             self._write_fut = self._loop._proactor.send(self._sock, data)
403             if not self._write_fut.done():
404                 assert self._pending_write == 0
405                 self._pending_write = len(data)
406                 self._write_fut.add_done_callback(self._loop_writing)
407                 self._maybe_pause_protocol()
```

The proactor object is again an instance of `loopProactor` defined in `asyncio/windows_events.py`. Inside the proactor `send` method, the function `WSASend` is invoked on line 537. This invokes the [Windows API](#), writing the bytes in the buffer on the socket via a Python C extension module

https://github.com/python/cpython/blob/main/Lib/asyncio/windows_events.py#L537

```
535         ov = _overlapped.Overlapped(NULL)
536         if isinstance(conn, socket.socket):
537             ov.WSASend(conn.fileno(), buf, flags)
538         else:
```

When the `send` operation completes, the transport write operation returns. Since the response is not complete (the body is yet to be sent), the conditional on line 546 evaluates to false, and the outer `send` method returns as well

Back in `starlette/responses.py` on line 171, the `send` function is invoked a second time with the HTTP response body

<https://github.com/encode/starlette/blob/master/starlette/responses.py#L171>

```
169         }
170     )
171     await send({"type": "http.response.body", "body": self.body})
172
```

This message is again passed up the middleware to the `send()` method defined in `uvicorn/protocols/http/h11_impl.py`. However, because the request has started, the conditional on line 481 evaluates to false this time, so the branch on line 515 is taken. This branch passes the bytes into an `h11.Data` object on line 529, which is used for abstracting

HTTP response bodies

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L529

```
528         else:
529             event = h11.Data(data=body)
530             output = self.conn.send(event)
531             self.transport.write(output)
```

This Data object is then passed to the send method on the connection object on line 530. This method from `h11._connection.py` converts the HTTP event into bytes that can be sent on the TCP connection

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L530

```
528         else:
529             event = h11.Data(data=body)
530             output = self.conn.send(event)
531             self.transport.write(output)
```

The output bytes are then passed to the write method on the transport object on line 531, which will send the raw bytes over the TCP socket

https://github.com/encode/uvicorn/blob/master/uvicorn/protocols/http/h11_impl.py#L531

```
528         else:
529             event = h11.Data(data=body)
530             output = self.conn.send(event)
531             self.transport.write(output)
```

The transport object is an instance of `_ProactorBaseWritePipeTransport`. The write method of this class is defined on line 338 of `asyncio/proactor_events.py`. This time, however, the transport is in the writing state. Thus, the conditional on line 363 evaluates to false, and it instead takes the branch on line 367. Critically, this updates the value of `_buffer` to be a mutable byte buffer containing the HTTP body bytes on line 369

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L369

```
367         elif not self._buffer: # WRITING -> BACKED UP
368             # Make a mutable copy which we can extend.
369             self._buffer = bytearray(data)
370             self._maybe_pause_protocol()
```

When the first future completes, it reruns the `_loop_writing` method in `asyncio/proactor_events.py`. This time, data is none, so line 388 assigns data to be the bytes stored in `_buffer` in the previous step

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L388

```

387         if data is None:
388             data = self._buffer
389             self._buffer = None

```

Then, on line 402, this data is sent using the proactor send method

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L402

```

401         else:
402             self._write_fut = self._loop._proactor.send(self._sock, data)
403             if not self._write_fut.done():
404                 assert self._pending_write == 0
405                 self._pending_write = len(data)
406                 self._write_fut.add_done_callback(self._loop_writing)
407                 self._maybe_pause_protocol()

```

This method sends the body data using the [Windows API](#). Specifically, it writes the supplied buffer to the socket using the WSAsend method

https://github.com/python/cpython/blob/main/Lib/asyncio/windows_events.py#L537

```

535         ov = _overlapped.Overlapped(NULL)
536         if isinstance(conn, socket.socket):
537             ov.WSASend(conn.fileno(), buf, flags)
538         else:

```

Note that the `_write_fut` again calls `_loop_writing` when this write completes. However, after writing the body data, `data` will be `None` and `_buffer` will be `None`. Thus, the conditional on line 390 will evaluate to `True`, and `_call_connection_lost()` will be scheduled in the event loop on line 392

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L392

```

390         if not data:
391             if self._closing:
392                 self._loop.call_soon(self._call_connection_lost, None)
393             if self._eof_written:
394                 self._sock.shutdown(socket.SHUT_WR)

```

In `_call_connection_lost()`, `asyncio` first notifies `uvicorn` that the TCP connection is being terminated by invoking the `connection_lost()` method of the protocol on line 158

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L158

```

157         try:
158             self._protocol.connection_lost(exc)
159         finally:
160             # XXX If there is a pending overlapped read on the other

```

Finally, on line 165, the socket is shut down. This is followed by the socket being closed on line 166, which terminates the TCP connection

https://github.com/python/cpython/blob/main/Lib/asyncio/proactor_events.py#L166

```
164         if hasattr(self._sock, 'shutdown') and self._sock.fileno() != -1:
165             self._sock.shutdown(socket.SHUT_RDWR)
166         self._sock.close()
167         self._sock = None
```

At this point, all of the bytes for the HTTP status line, headers, and body for the response have been sent over the TCP socket to the client. Additionally, the TCP connection has been closed gracefully. Thus, the client has exactly what it needs to display the web content

These technologies are useful because they abstract many complexities of building a web application. Specifically, they completely abstract the process of managing an event loop as well as interfacing with the operating system to create a TCP listening socket, accept connections, read data, and send data. They also abstract the process of parsing bytes and managing middleware. All the application programmer needs to do is follow the framework, and they can build functional applications in very little time (hence the name FastAPI 😊)